

Agiwal, Ankur, et al. "Napa: Powering scalable data warehousing with robust query performance at Google." *Proceedings of the VLDB Endowment* 14.12 (2021): 2986-2997.

Google Napa

Abstract

- ▷ Napa: Analytical data management system
 - ▷ principle technology: "aggressive use of materialized views which are maintained consistently as new data is ingested"
- ▷ Design principles: ① Robust Query performance
 - ② Flexible config. of client DBs
- ▷ Main Challenge: Napa had to meet existing requirements
 - ▷ - not a greenfield system

Introduction

- ▷ High level constraints
 - sub second query performance
 - petabytes of data
 - continuously updated
 - consistent & fresh query results
 - fault tolerant
- ▷ Began from set of requirements
 - less flexibility / not starting from clean slate
- ▷ Built to replace Mesa
 - historical data onboarded + new clients added
 - "Broader mandate" than Napa
- ▷ "Bedrock" principles of design
 - ① Robust Query Performance
 - low latency + low variance in latency regardless of load
 - Napa guarantees robust query perf. w/ consistent results
 - ② Flexibility
 - not all clients have same perf. reqs. or are willing to pay
 - ③ High throughput Data Ingestion
 - implements distr table & row maintenance framework based on Log Structured Merge Tree

- ↳ current flexibility (what) ...
- Materialized views are updated consistently
 - ↳ Contrast to current trends of other systems which achieve perf. via efficient scans of base data
 - Tradeoffs
 - coverage affects query performance
 - refresh rate affects freshness
 - clients can specify both based on their needs

Design constraints

- 3 objectives: query performance, freshness, computational cost
 - Q perf: measured by latency
 - freshness: measured by time between row being added to table & being available for querying (ranges from few minutes - couple hours)
 - cost: ingestion & maintenance dominate
- Q perf must also consider predictability
- Clients must make three way tradeoff, all have diff needs
- Coupling ingestion w/ storage: new data can't be ingested until fully processed
- Coupling new data w/ querying \Rightarrow slows down perf
- Neither approaches work b/c
 - 1) freshness is sacrificed: ingestion dependent on storage bandwidth
 - ↳ stale data is served
 - 2) Q performance consistency is sacrificed: views could be gen on the fly (lazy), but Q perf is inconsistent (fresh, but slower)

Design Choices

- Ingestion \rightarrow Storage \rightarrow Query/ Serving
- Ingestion f.w. is responsible for committing updates (called deltas) into tables. Write optimized
- Storage f.w. incrementally applies updates to tables & views
 - tables & views maintained as LSM forest
 - a table is a collection of updates

- compaction: deltas consolidated to form larger deltas

- table deltas get transformed into view deltas

3) Query Serving: answers client queries

- system performs merging of table/view deltas @ query time

- Q latency dependent on Q time merging

- ↳ faster storage can proc updates, fewer deltas need to be merged @ Q time

- Ingestion decoupled from view maintenance & view maintenance from Q proc.

- ↳ enables client to fine tune to their needs

- Low Effort: cost optimized \Rightarrow reduced Q perf

- ↳ translates to less aggressive compactions & more merging @ exe time

- can also be less materialized views or reduced freshness

- High Effort: better Q perf

- ↳ low fan in merges @ Q time

► Providing Flexibility

- Users can specify needs which get translated into system configs

- ↳ # views, max # deltas that can be opened during proc

- System isn't static due to ingestion so Napa offers Queryable

- Timestamp: live marker that indicates freshness

- Now - QT = data delay

- All data up to QT timestamp can be queried by client

- Guarantee that all data used to serve Qs will meet conditions for delivering expected Q perf

- Categories of clients:

1) Freshness tradeoff

- client needs good Q perf w/ moderate cost

- Napa maintains moderate # views & fewer delta merges @ exe

- exe fw uses fewer workers & "cheap opportunistic resources"

2) Q perf. tradeoff

- client needs fresh data, but has lower perf standards

- system maintains less views, but more deltas can be merged
- more workers alloc for ingestion than view maintenance

3) Cost tradeoff

- client needs good perf & freshness
- low # deltas merged @ run time + more views are maintained
- many workers allocated to ensure faster ingestion, high throughput view maintenance

• Data Availability

- Napa guarantees availability @ scale of data center outages
- client DBs replicated @ multiple datacenters
- the consistency across datacenters
 - decouples data & metadata ops
 - one approach: exe ingestion activity as synchronous transactions
 - Napa approach: decouple exe of data & mdata ops
 - ↳ data ops are synchronously, mdata ops used periodically to ensure replicas remain synchronized
 - BT indicates time @ which all tables & views are consistent across all data centers

• System Architecture

- high level arch consists of data & control planes
 - data: ingestion, storage query serving
 - control: controller that coord. work among various sub systems
 - also responsible for sync & coord. mdata ops across datacenters
- built by leveraging existing Google infrastructure components
 - Colossus: table in Napa is collection of Colossus files
 - Spanner: functions that req specific transaction semantics (mdata)
 - F1 Query: Query serving + large scale data proc like view creation & maintenance
- Clients use ETL pipelines to insert data into tables
 - delivered to any replica & Napa incorporates to all datacenters



- ▷ • Napa excels @ workloads w/ aggregations & complex filters (dashboards)
- ▷ ↳ storage & view maintenance is key to maintain these aggregations
- ▷ • low latency achieved by directing Qs to materialized views
- ▷ & not base tables
- ▷ • low variation achieved by controlling fan in merges
- ▷ • Napa tables & views are sorted, indexed, & range partitioned
- ▷ ↳ motivated by strict latency & resource reqs which necessitate
- ▷ indexed key lookups
- ▷ • addl partitioning schemes being investigated
- ▷ • Napa relies heavily on merging & sorting performance
- ▷ ↳ sorting, merging, & GB were optimized
- ▷ • Controller schedules compaction & view updates to keep
- ▷ count of deltas such that it can be configured
- ▷ • most clients req freshness of 10s of minutes
- ▷ • System has been able to maintain all client reqs wrt freshness

▶ Ingestion

- ▷ • goal of i.f.w. is to allow pipelines to insert w/o much overhead
- ▷ • seeks to accept w/ minimal proc & make data durable w/o being
- ▷ concerned w/ view maintenance
- ▷ • ingested rows get mdata timestamp for ordering & get marked as
- ▷ committed after meeting durability reqs
- ▷ • clients can specify # tasks dedicated to ingestion

▶ Queryable Timestamp: indicates freshness of data

- ▷ • only advances when ingested data meets performance requirements
- ▷ • clients don't see any data after QT
- ▷ • clients can easily use QT to tune performance configs
- ▷ • Good Q perf: optimize for reads + ensure views available
- ▷ • Table is collection of all of its delta files
- ▷ • non-queryable deltas: new deltas (typically few seconds)

- deltas correspond to updates over window of time
 - ↳ largest span weeks or months
- Ea delta sorted by key, range partitioned, has "B-tree like index"
- deltas merged as needed @ query time
- Underlying data is a column store, but has to manage m-views
 - & fast lookup, so it borrows ideas from row stores like B-Trees
 - & PAX layouts in physical design
- Q latency limits max # deltas that can be merged (@ a time)
- BT is the delta that bounds this #
 - usually 10s of deltas
 - limit is auto adjusted
 - high perf req \Rightarrow lower # & vice versa
 - As # inc there is a tail effect
 - by keeping near constant, Map guarantees variance in Q latency is low
 - Q dep on background ops
 - there is a DB level BT (min of all table BTs)
 - ea replica has a local BT which constitute the global BT
 - Map will use replicas that have a min BT based on replica availability requirements (will direct Qs to certain replicas)

■ Maintaining views at scale

- storage sub system is resp for maintaining views & compacting deltas
- also resp for ensuring data integrity, durability, & handling outages
- skew in view maintenance
 - occurs in process of transforming updates on base tables into updates on the views
 - mapping of base key space into view key space may lead to discontinuities may map to a narrow view key range
 - system must auto adjust to skew to ensure that BT is not susceptible to straggler views or tables

- storage sub system adjusts to cost budgets

- Key aspects of view maintenance

- 1) F_1 Query as relational data pump to compact tables & maintain views

- 2) Replanning to avoid skew

- system can replan on the fly if it detects data skews

- 3) Intelligence in the loop (for tail mitigation)

- data centers are specifically chosen for task exe based on historical load, active straggler task termination, concurrent task exe

- ▶ Query Optimizations challenges in view maintenance

- view maintenance system exploits data properties in the input

- ↳ sorted, partitioned (hard to recreate)

- Ext data property: sort order of view relative to base table

- one approach: resort view keys based on view sort order regardless of base table sort order \Rightarrow @ scale: expensive

- instead beneficial to preserve input sortedness as much as possible even if views sort order & base table sort order only partially overlaps

- 3 classes of views (based on cost of maintenance)

- 1) cheapest: those that share prefix w/ base table (ABC, AB)

- 2) partial prefix w/ base table (ABCD, ABD)

- cluster on AB, sort on D

- 3) most exp: no shared prefix (ABCD, DC A)

- req repartitioning & re-sorting

- for views w/ high cardinality reduction:

- preserving sort order isn't as important (resort if needed)

- w/ low cardinality reduction (similar size as table)

- sort & merge becomes imp

- Napa employs "state of the art" sort library

► Mechanics of compaction

- compaction combines mult. input deltas into single output delta
- improves Q perf & reduces storage consumption by
 - 1) sorting input rows together
 - 2) aggregating mult updates to same rows
- Compacting is expensive for high ingestion tables
 - reduce data freshness by delaying when data becomes queryable
- compaction is merge sorting b/c delta files are sorted individually
- fan in kept large to minimize tree height
- up to 1000 inputs before merge perf dec.

► Robust query serving performance

- fast Q results still req even for large workloads
- Reducing Data in the critical path
 - uses views whenever possible instd of base table
 - pushes down filters & partial aggregations to min transfer to F1
 - ↳ Napa storage & F1 & workers not always in same data center
 - maintains sparse Btree indexes on stored data & uses this to partition input Qs into thousands of subqueries

► Min. Sequential I/Os

- common to hit high latency if data has to be read from disk
- Napa uses Q to det which mdata to use which det what data is processed. Napa ensures all mdata served from mem
- All data reads go through transparent distr data caching layer
 - ↳ shares work of subqueries to ensure overlapping reads are only processed once
- Online & offline prefetching are performed to further reduce I/Os
 - offline: newly ingested data is cached for frequently Q'd tables
 - Online: "Shadow Q executor" runs ahead of main executor by skipping proc. & shares data access pattern
 - ↳ more accurate prefetching

► Combining small I/Os

- Napa parallelizes Qs I/O calls across deltas & Q2 columns
- However, high parallelism can lead to tail latency
- Napa uses BT to limit # deltas & combines small I/Os by

1) Lazy merging across deltas

2) Size Based Disk Layout

- Napa supports multiple disk layout options based on delta size
- PAX - small deltas: good for lookup queries
- column by column - large deltas:
 - efficient for scans, but req 1 I/O per column for lookup
- columnar benefits w/ reduction in I/O ops

► Tolerating Tails & Failures

- focus on tolerating tail latency rather than eliminating it
- Non-Streaming MPC:
 - hedging: secondary MPC sent after a given delay & waits for faster reply
- Streaming: Napa est its progress rate & req server exiting to report progress - if it falls behind, new streaming MPC is started on diff server w/o losing progress
- For data center wide issues, Qs are routed to other data centers

■ Production Metric Insights

- billions Qs a day & trillions of rows ingested

► Views & BT help achieve robust & perf

- reading from views reduces variance & improves Q perf + tail latency
- plateau @ 8 views as this is point where base tables no longer needed
- reducing # deltas @ a time
 - 1) less small parallel I/Os which cause tail latency
 - 2) data is premerged (compacted)

► Handling Infrastructure Issues (see figure 9)

- Q perf remains stable despite ingestion changes or outages

► Client Workloads (previously discussed)

- see 9.3

► Related Work

- Napa provides continuous ingestion w/ high perf & tunable freshness
- further adv idea of disaggregation
- Tradeoffs btwn freshness or higher cost (tunable params)
- Napa is fully indexed sys optimized for Key lookups, range scans, incremental maintenance of indexes on tables & views
- supports wide variety of ops
- drop in replacement to Mosa
- ↳ sig improved @ latency & costs to run
- Napa use B tree variant LSM tree, which are optimized for updates
- LSM trades high write throughput for fast reads
- ↳ deltas solve this

► Conclusion

- serves dashboards, apps, internal users
- similarities to existing systems (available, scalable, etc)
- unique: materialized views, maintaining views while ingesting, tunable
- Views optimized for continuous high low insertions using LSM forests
- BT is live indicator of performance
- serves clients w/ diff needs

Poppe, O., Guo, Q., Lang, W., Arora, P., Oslake, M., Xu, S., & Kalhan, A. (2022, January). Moneyball: Proactive Auto-Scaling in Microsoft Azure SQL Database Serverless. VLDB 2022, 1279–1287.
Retrieved from <https://www.microsoft.com/en-us/research/publication/moneyball-proactive-auto-scaling-in-microsoft-azure-sql-database-serverless>

Moneyball

Abstract

- ▷ Serverless compute currently auto scales based on workload demand, but only reactively
- ▷ ↳ resources may not be immediately available when customer comes back online after idle period
- ▷ • Goal is to predict pause/resume patterns to reduce delay + avoid taking away resources for short idle periods

Intro

- ▷ • customers charged per second
- ▷ • serverless can have delays, so for time sensitive apps. provisioned compute may be more suitable
- ▷ • goal is to proactively resume based on historical resume patterns + avoid freezing resources for short pauses

Challenges

- ▷ 1) large search space of tunable params
 - ▷ - too many to search exhaustively
 - ▷ - identify trends of how params affect results to choose reasonable set
 - ▷ - param choice involves tradeoff b/wn QoS & COGS
- ▷ 2) Opposing optimization objectives
 - ▷ - 2 goals (proactive resumes & avoid short pauses) are opposing b/c incorrectly predicted resumes will inc short pauses
 - ▷ - reducing short pauses will make predictions harder b/c fewer hist. resumes
- ▷ 3) Changed resource usage patterns
 - ▷ - "negligible" % of provisioned or serverless DBs follow strict pattern
 - ▷ - provisioned is short lived & underutilized, serverless is opposite
- ▷ • Moneyball seeks to find middle ground b/wn proactive resume & avoiding short pauses
- ▷ • Make recommendations to tune key params + reduce costs of incorrect resumes until they are used

• To reduce short pauses, 2 solutions considered:

- 1) put cap on # pauses per DB/day
- 2) wait time before scale down (greedy & predictive considered)

• Introduces visual component to compare effect on AOS & COGS

• transfer lessons learned from provisioned

• results span half year in 10s of Azure regions \Rightarrow generalizable

• Up to 80% proactive resumes are correct & 99% long lived DBs benefit

• avoid short pauses w/ logical pause (wait time). Avoids up to half of pauses & 49% of DBs benefit

Moneyball problem

• Provisioned results in underutilization & extensive idle periods unless manual scale down occurs

• Manual scale down is labor intensive, expensive, error prone

• If online, serverless scaling usually has low latency b/c some resources remain alloc.

• If paused latency increases -- larger performance hit

• Current Approach is to resume/pause on login/logout

• Correct Proactive Resume: correct if resources of a DB are used w/in a window

• aim to min operational costs assoc w/ proactive resume

• short pauses are ineffective b/c resources can effectively be reused

• Ineffective: ineffective if pause falls w/in given duration

• If pauses do not occur for long idle periods, resources are wasted

• Goals: ① Max # correct proactive resumes ② Min # short pauses

③ Min op costs of these 2 optimizations

Transfer learning from provisioned to serverless compute

Features

• half year, 10s of regions

• variables: timestamp, DB ID, DB state, duration resumed/paused, DB compute capacity (max vCores), DB creation & deletion timestamps, region

► ML models

- Provisioned used various ML models to predict load: ARIMA, Prophet, NimbusML, Neural nets, Exp. Smoothing, Persistent Forecast heuristic
 - used given day to predict following day
 - NimbusML was most accurate, but not much more than persist. forecast.
 - most DB fall into one of 2 categories:
 - Easily predictable (stable load or predictable pattern)
 - Hard to predict (load is random)
- Serverless grouped into categories
 - stable: resumed or paused @ least x% of the time in historical data
 - pattern: x% of pauses & resumes occur w/in a given window over time
 - predictable: stable or pattern
 - most serverless DBs are predictable
 - 74% stable @ least 90% of time
 - 3% follow pattern w/in 15 min @ least 90% of time
 - 23% are unpredictable
 - trained ML models on 3 weeks of data & pred a given day
 - measure hinge loss
 - NimbusML was again most accurate

► Proactive Resume

► Algorithms

- Example DB is unpredictable, but there is a predictable recurring resume
- prob resume rec. = $\frac{\# \text{ weekdays resumed in window}}{\# \text{ weekdays in window}}$
- recommended to proactively resume DB if prob is \geq threshold (Algo 1)
- time complexity: $O(\# \text{ DBs in set} \cdot \# \text{ windows/day} \cdot \# \text{ rows in hist data}^*)$
 - * in the given window & set of DBs
- recommend to proactively resume if there @ start of window on day if there exists a resume in the predicted pattern

- same as algo 1, except has to predict patterns
- time complexity = $O(\# \text{ DBs} \cdot (\text{predict function} + \# \text{ windows} - \# \text{ predicted patterns}))$

► Middle Ground between DBs & COGS

• total cost savings = $\sum_{s \in S} \text{pauses}(s) \cdot \text{vcores}(s) \cdot \text{cost}$

• wasted cost = $\sum_{s \in S} \text{wait}(s) \cdot \text{vcores}(s) \cdot \text{cost}$

- resume cost index: ratio of wasted costs to total cost savings

↳ depends on several tunable params (size window, length hist params)

- Long lived DB: exists for 2-3 weeks

- size of window: varied from 1-9 hours, measure % correct & wrong proactive resumes

- % correct inc from 22 to 56 for probabilistic as window grows

- % correct inc from 63 to 80 for predictive

- Prob benefits 25 - 62 % of DBs & pred 99 %

- % correct resumes & benefited DB is up to 3x higher for pred than prob

- inc window results in inc cost b/c resources are idle longer

- Prob has up to 5x fewer wrong resumes \Rightarrow 5x lower cost idx

- Length of historical data: varied from 3-7 weeks

- 3 weeks: 36% correct resumes pro & correct, 43% DB have correct proactive resumes, 12% proactive & wrong

- percentages dec as length of hist data grows

- pred has up to 18x more wrong resumes \Rightarrow cost idx 10x higher

► Avoiding Ineffective Pauses

• short idle pauses ineffective + unnecessary workloads introduced

• avoid these by budgeting # pauses & delaying pauses

► Budgeting Algos

• one approach: restrict # pauses per DB/day & prioritize long pauses

• Greedy Budget: don't take length of pause into account & just allow first K pauses

- worst case: spends budget on early short pauses & avoids later long pauses

• Wasted cost = $\sum_{s \in S} \text{avoided}(s) \cdot \text{values}(s) \cdot \text{cost}$

- Predictive Budget: prioritizes pred. long pauses over pred short pauses when spending the budget

- algo takes predictions as param & gets longest pred pauses & avoids all pauses that don't start w/ a time delta of a long pause

- inc. overhead of making the predictions

- Optimal Budget: avoid top K shortest pauses per DB

- costs less than greedy by one order of magnitude

- Predictive runs the risk of making wrong preds which results in the available budget not being spent -- avoids up to 5x more pauses than greedy

- for some cases predictive is more exp & for others it is cheaper

- Budget can be defined @ different granularities & time windows

- Cost summary:

- Greedy order of mag higher than optimal
- pred can't guarantee lower cost than greedy

- Logical pause based algos:

- logical pause (logout time, resources taken away)
- physical pause (time resources taken away, log in time)

- Greedy logical pause:

- pauses DB for interval of time when customer logs out
- resources are still available in case customer returns
- reduces delays & short pauses, but resources are idle & long pauses get shortened \Rightarrow greedy is expensive

- wasted cost = $\sum_{s \in S} (\text{avoided}(s) + l \cdot \text{allowed}(s)) \cdot \text{values}(s) \cdot \text{cost}$

- Predictive Logical Pause:

- avoids predicted short pauses w/o reducing dur. of pred long pauses
- looks @ pred pauses & avoids those that are shorter than logical pause
- predicts pattern per DB per day

- Optimal logical pause: avoids all pauses shorter than logical pause
- # pauses & cost depend on length of pause
- greedy & optimal algos avoid 26-70% of pauses & benefit 33-58% of DBs as dur inc from 1-11 hrs
- greedy costs up to 6x more than original
- pred pauses tend to be shorter than actual & avoids up to 19% more pauses than greedy \Rightarrow cost is up to 4x higher
- most DBs benefit from greedy @ low cost

► Putting it all together

- Proactive Autoscale on serverless compute does good job @ summing interplay of ops
- In reactive, no proactive resume/pause so no COGS wasted, but delay inc
- introduced wait time interval to avoid ineffective pauses (logical pause)
 - wastes some COGS b/c resources are idle
- upto 90% of all resumes are proactive & correct w/in several hrs of long lived DBs & when combined w/ logical pause avoids up to half of pauses (Moneyball approach)

► Related Work

- Moneyball achieves contradictory goals of enabling proactive resume & avoiding short pauses while controlling operational costs

► Conclusion

- 2 optimizations introduced:
 - 1) reduce delays in resource availability by pred resume patterns per DB over time & proactively resume resources
 - 2) To reduce backend workload, avoid short pauses by logically pausing a DB that becomes idle
- results are used in all Azure regions